

NODEJS PERFORMANCE ANALYSIS

CHRISTIAN HENNIG

ScalaHacker @ inoio.de



caraboides



caraboides



DISCLAIMER

I'll show you what I've tried to solve my problem.

DISCLAIMER

I'll show you what I've tried to solve my problem.

But I failed! I did not find the solution for my problem.

DISCLAIMER

I'll show you what I've tried to solve my problem.

But I failed! I did not find the solution for my problem.

Maybe you have an idea. ;-)

OVERVIEW

OVERVIEW

- What's my problem

OVERVIEW

- What's my problem
- How to find bottlenecks with monitoring

OVERVIEW

- What's my problem
- How to find bottlenecks with monitoring
- How to get and analyze CPU profiles

OVERVIEW

- What's my problem
- How to find bottlenecks with monitoring
- How to get and analyze CPU profiles
- How to get and analyze heap dumps

OVERVIEW

- What's my problem
- How to find bottlenecks with monitoring
- How to get and analyze CPU profiles
- How to get and analyze heap dumps
- Tracing with DataDog

TASCHENHERZ

PROBLEM

PROBLEM

SERVICE RESPONSE TIME > 1SEC


PROBLEM

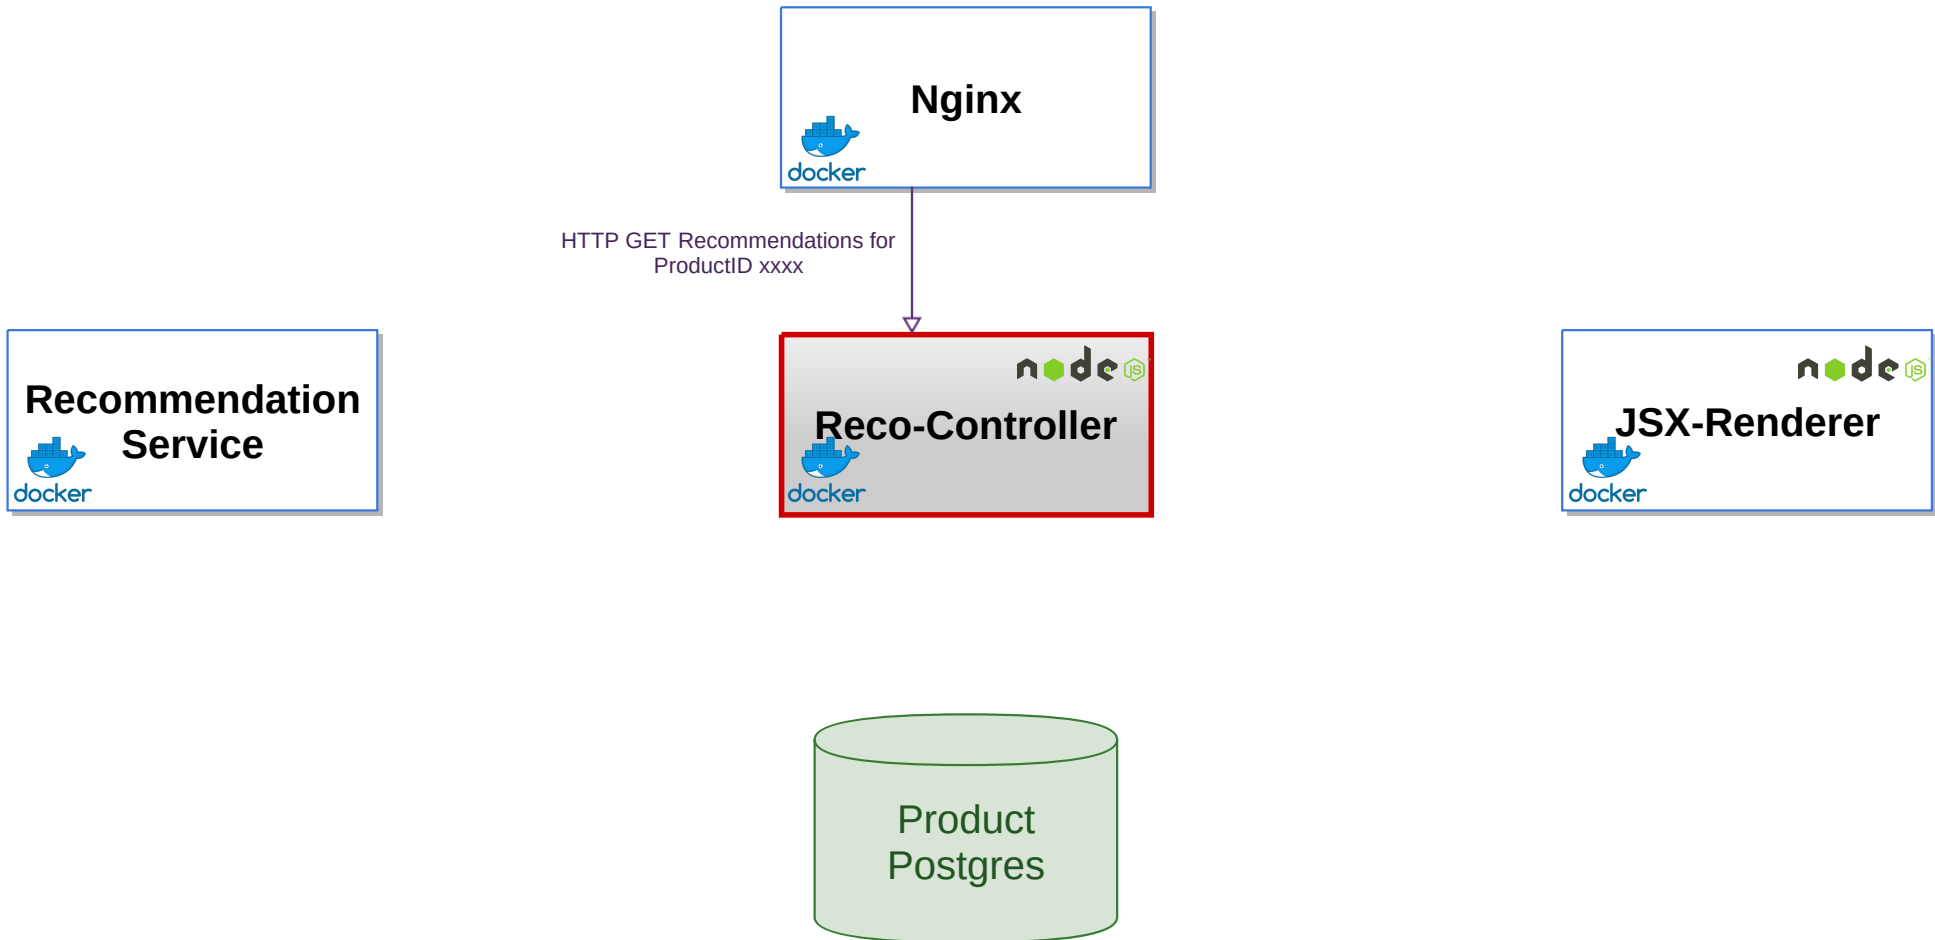
SERVICE RESPONSE TIME > 1SEC

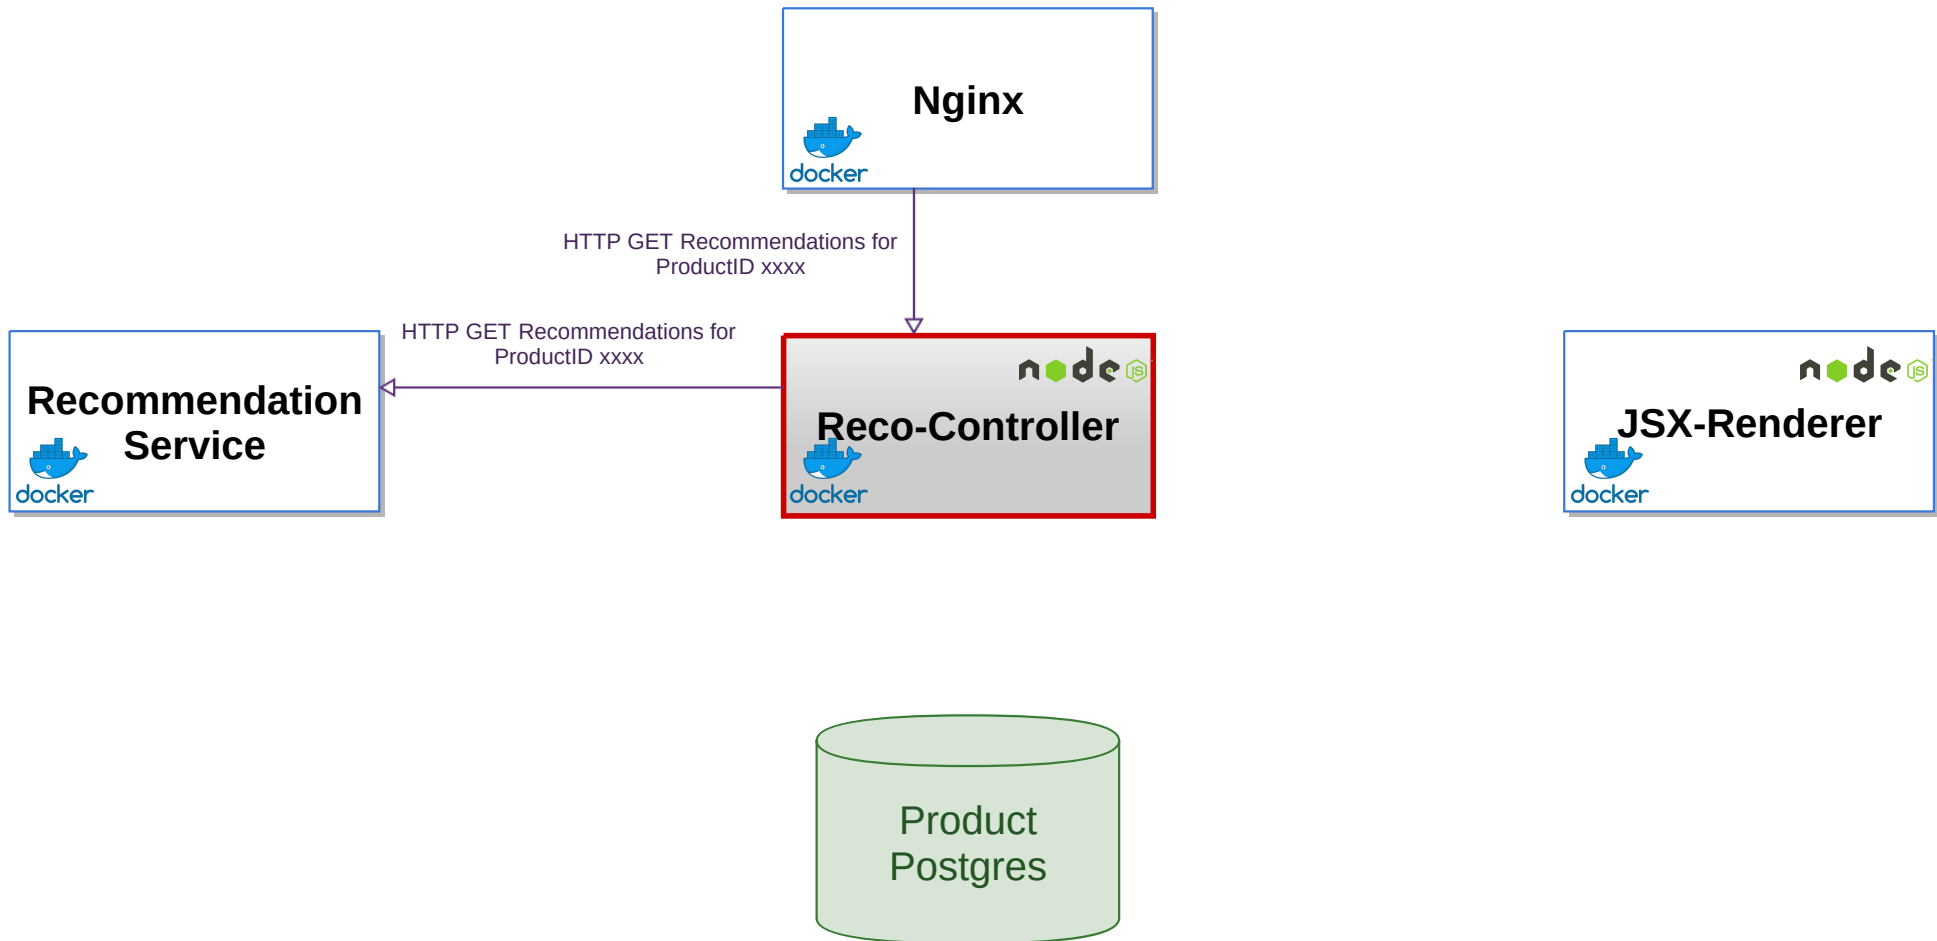
NGINX ABORTS REQUESTS AFTER 1 SEC

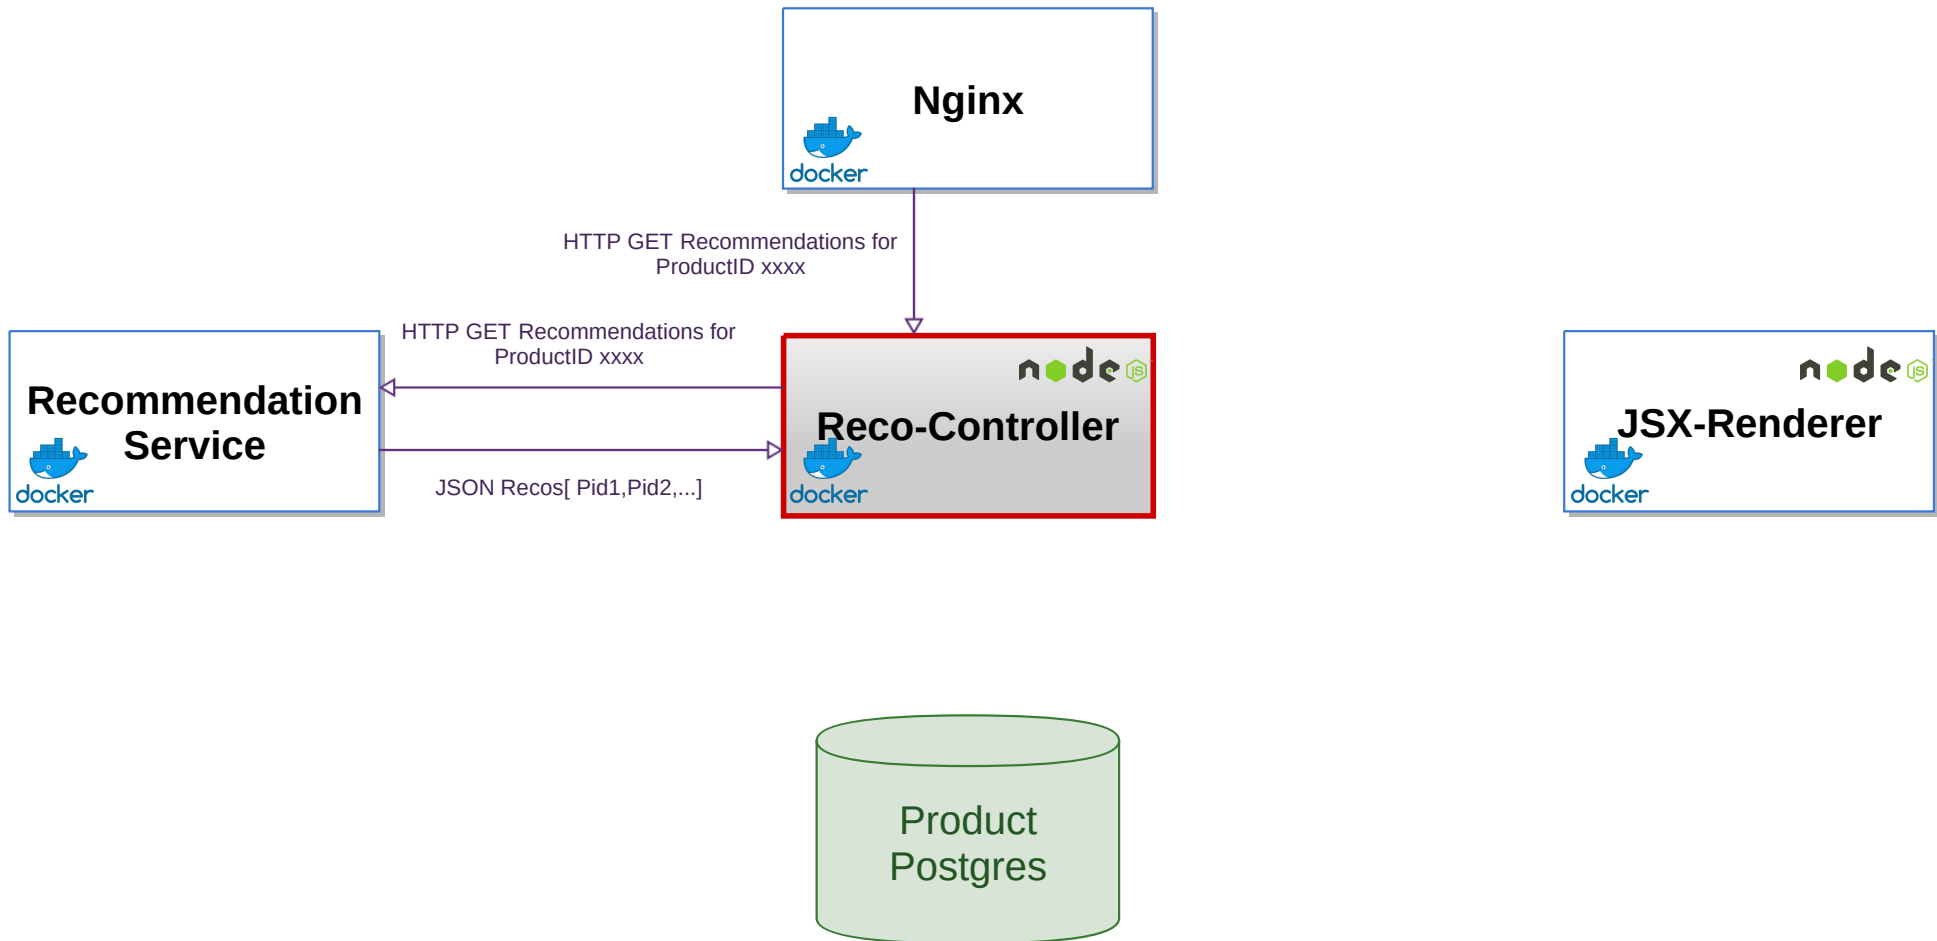


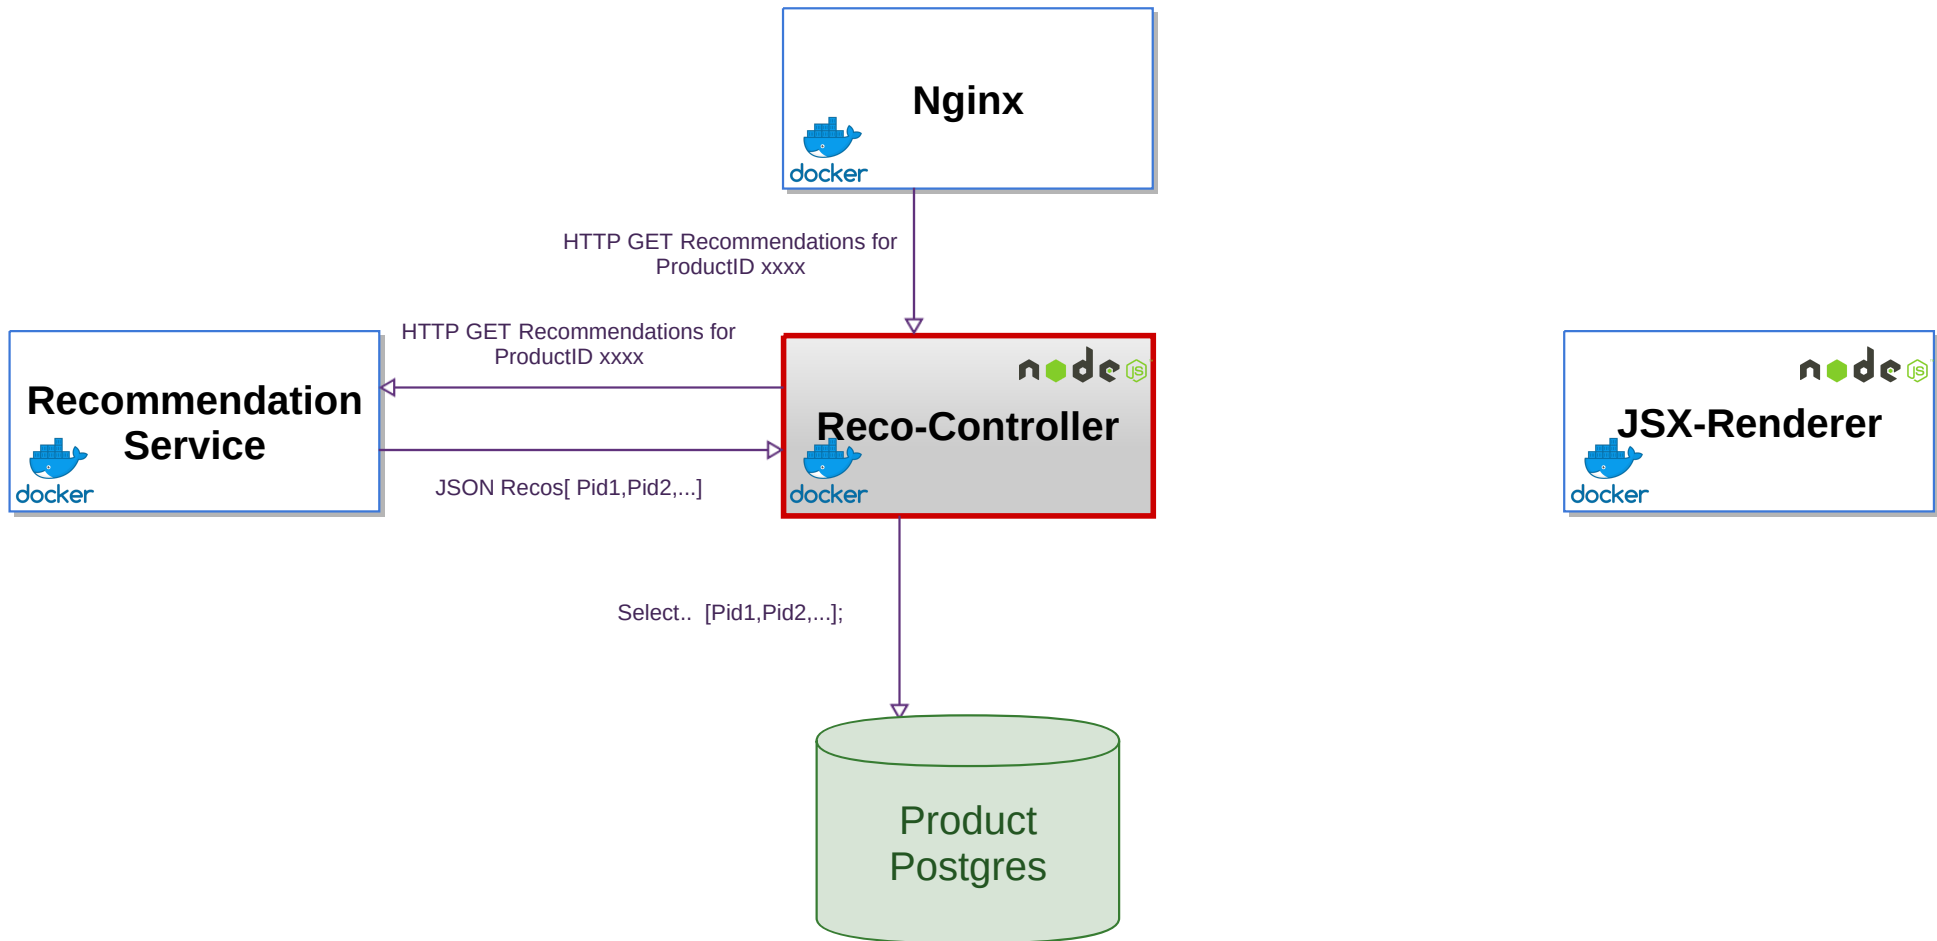
MicroServices:-)

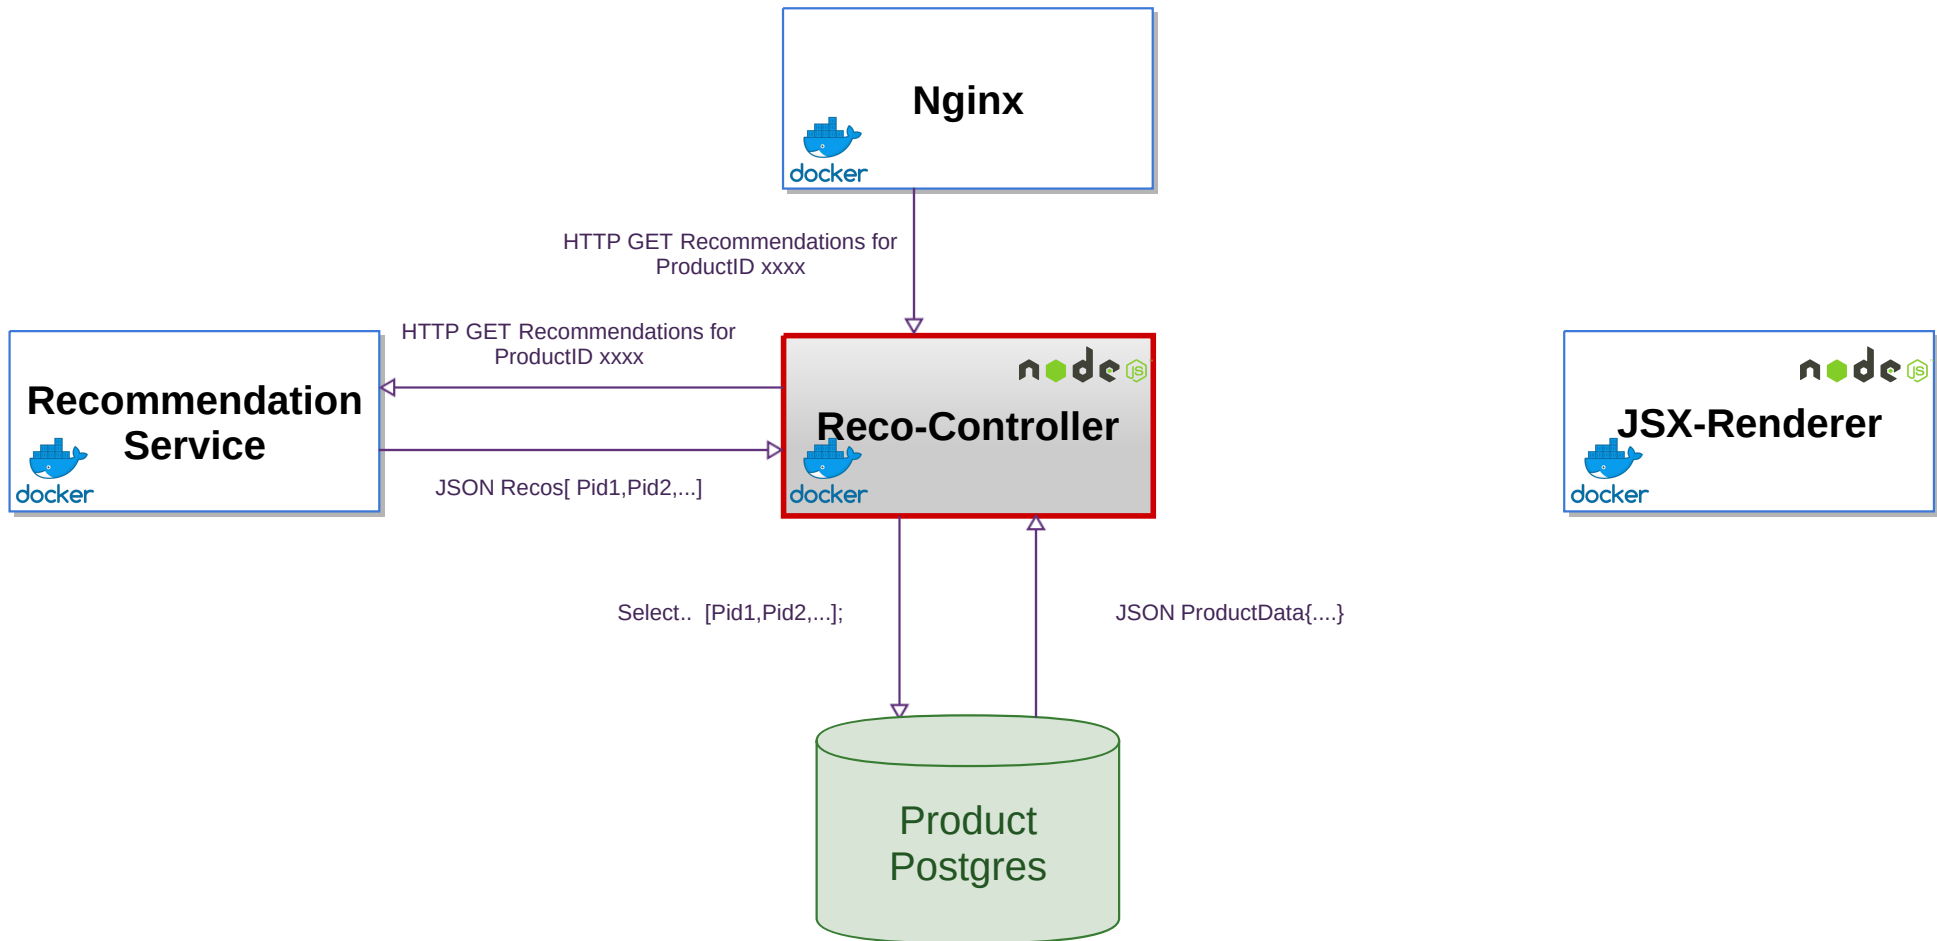


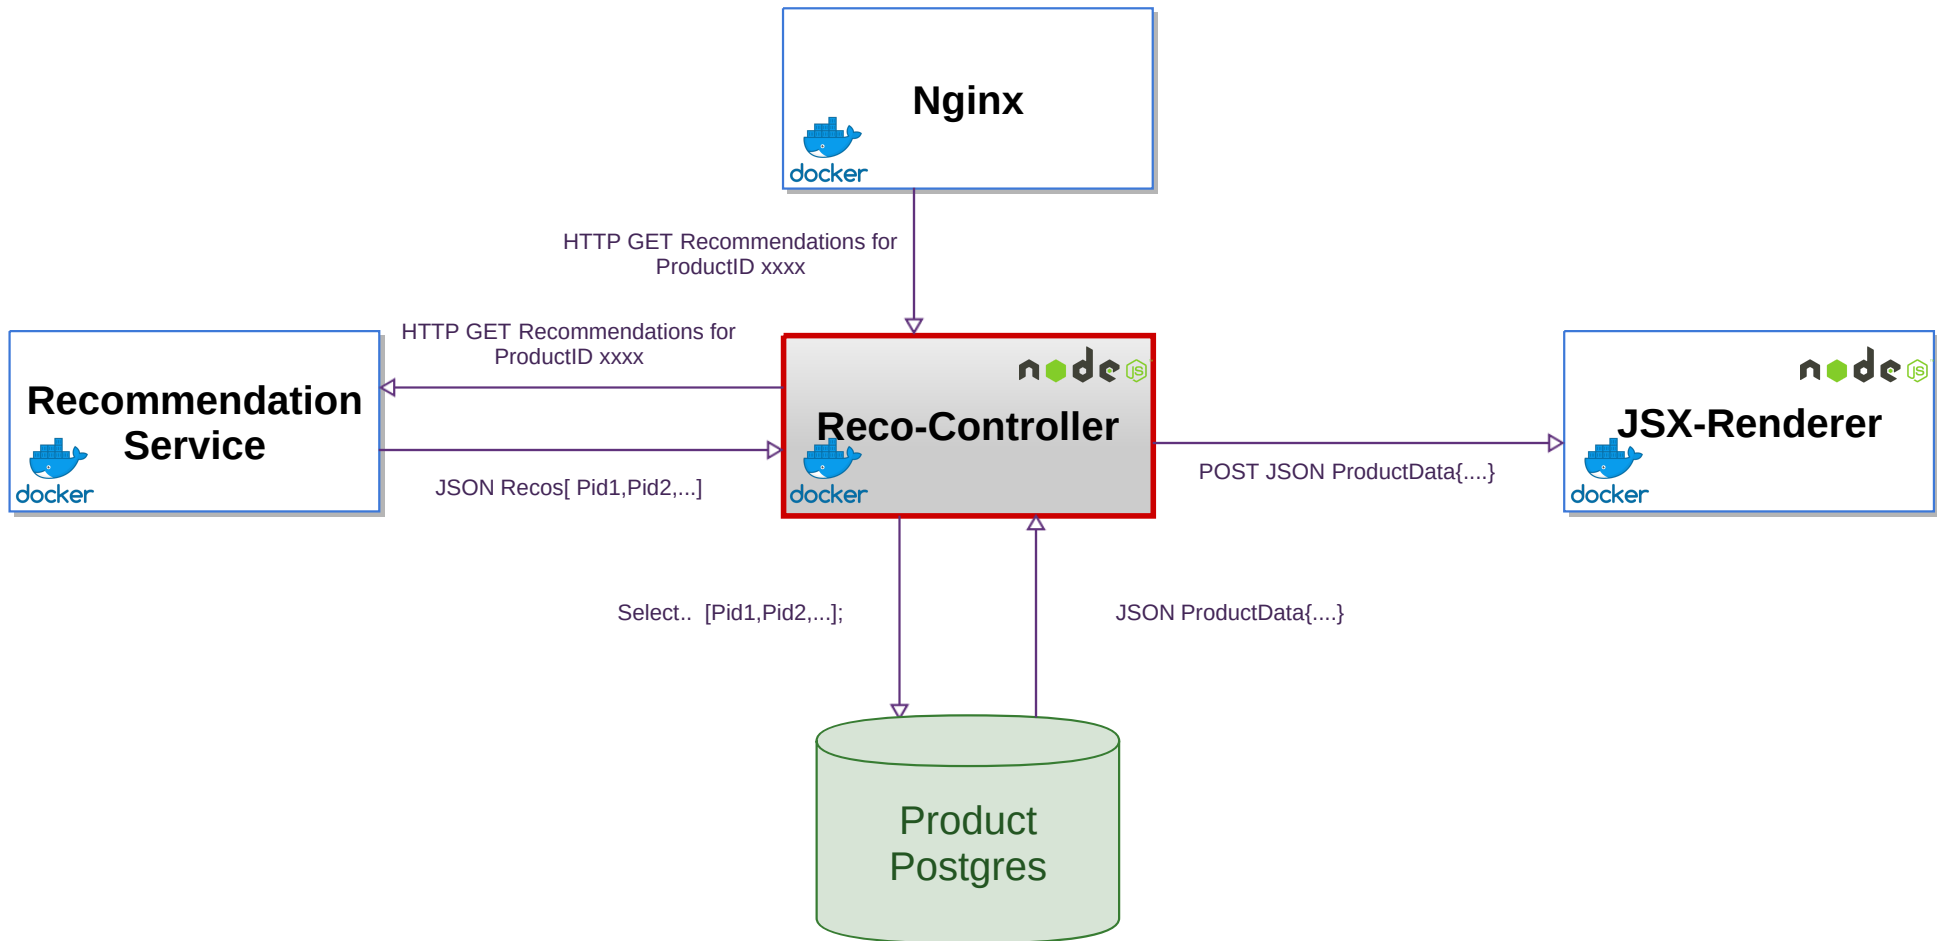


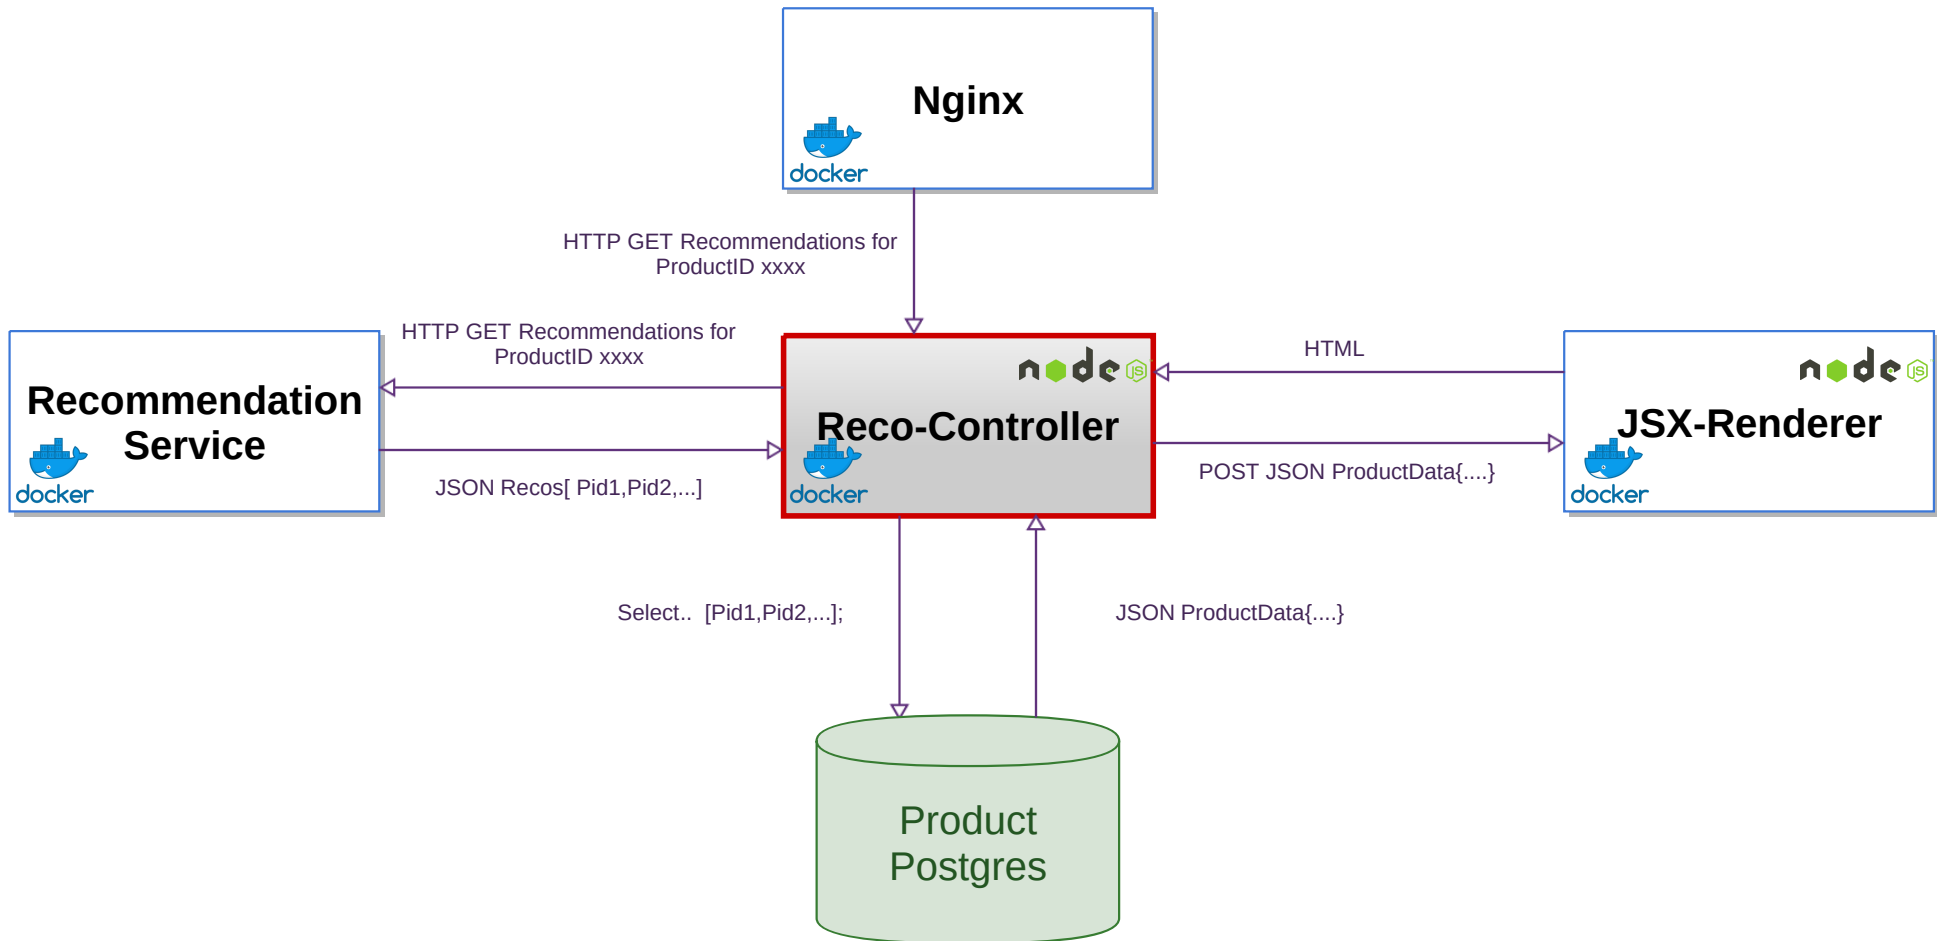


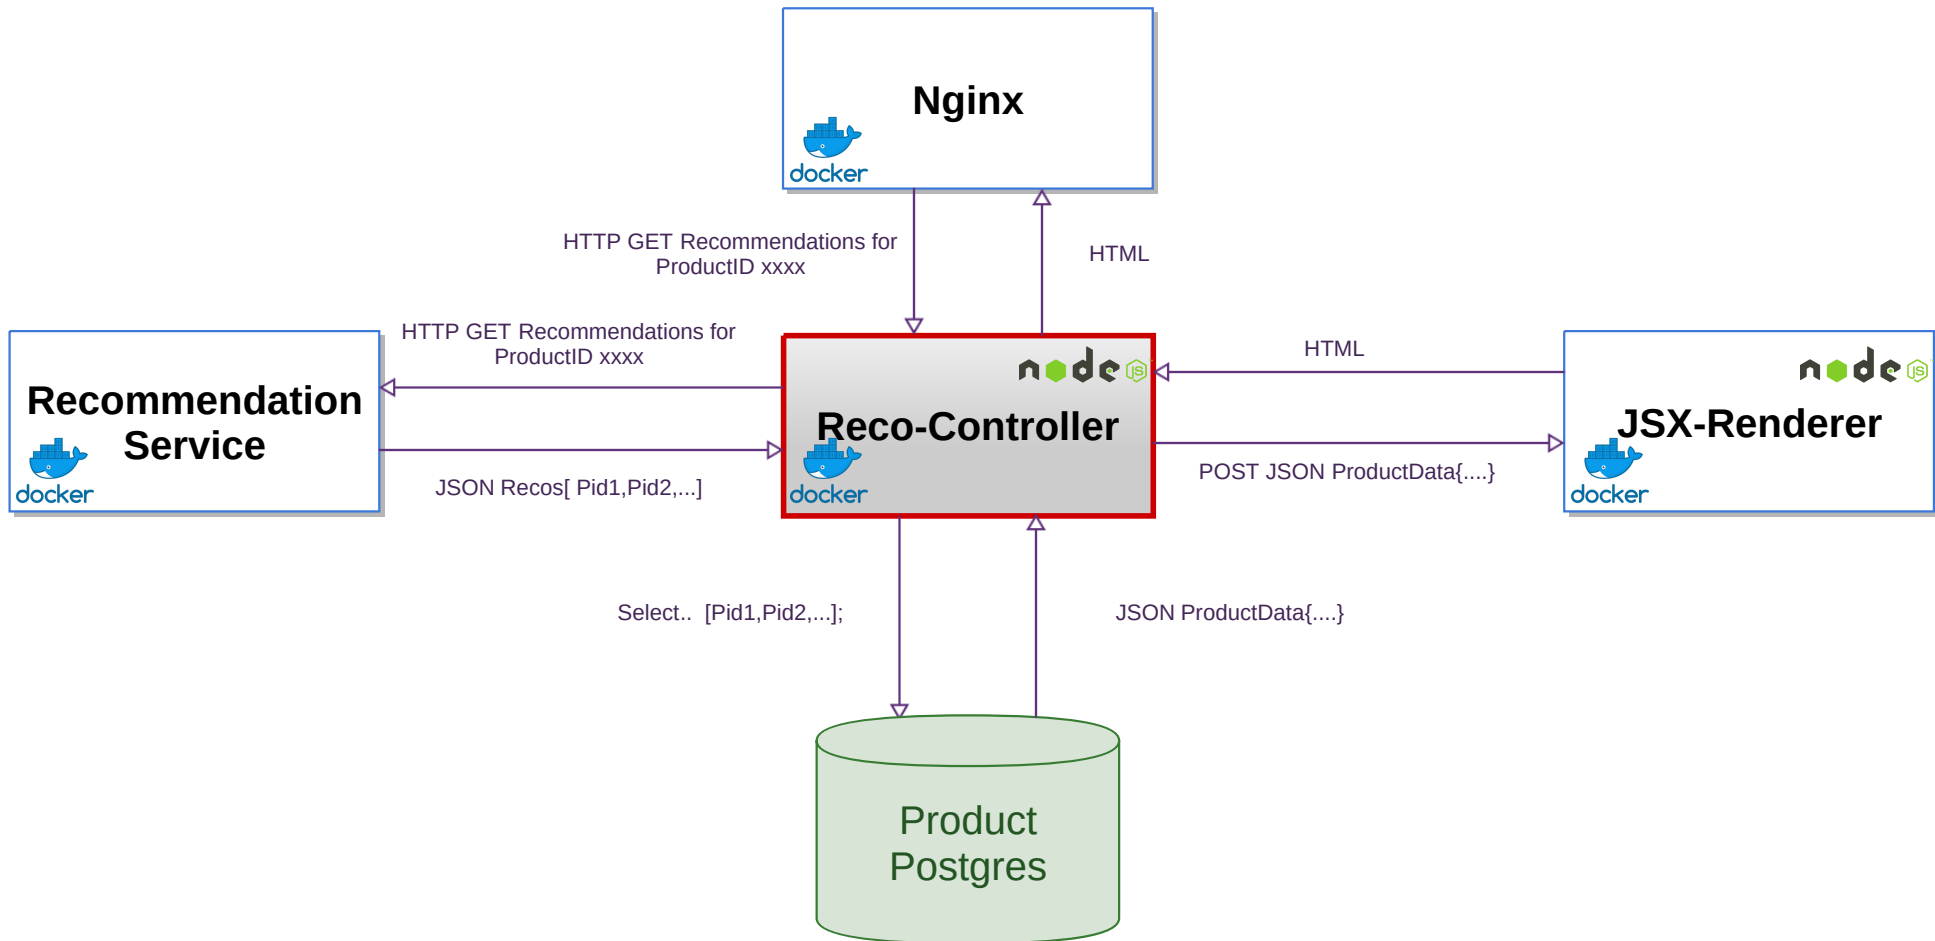


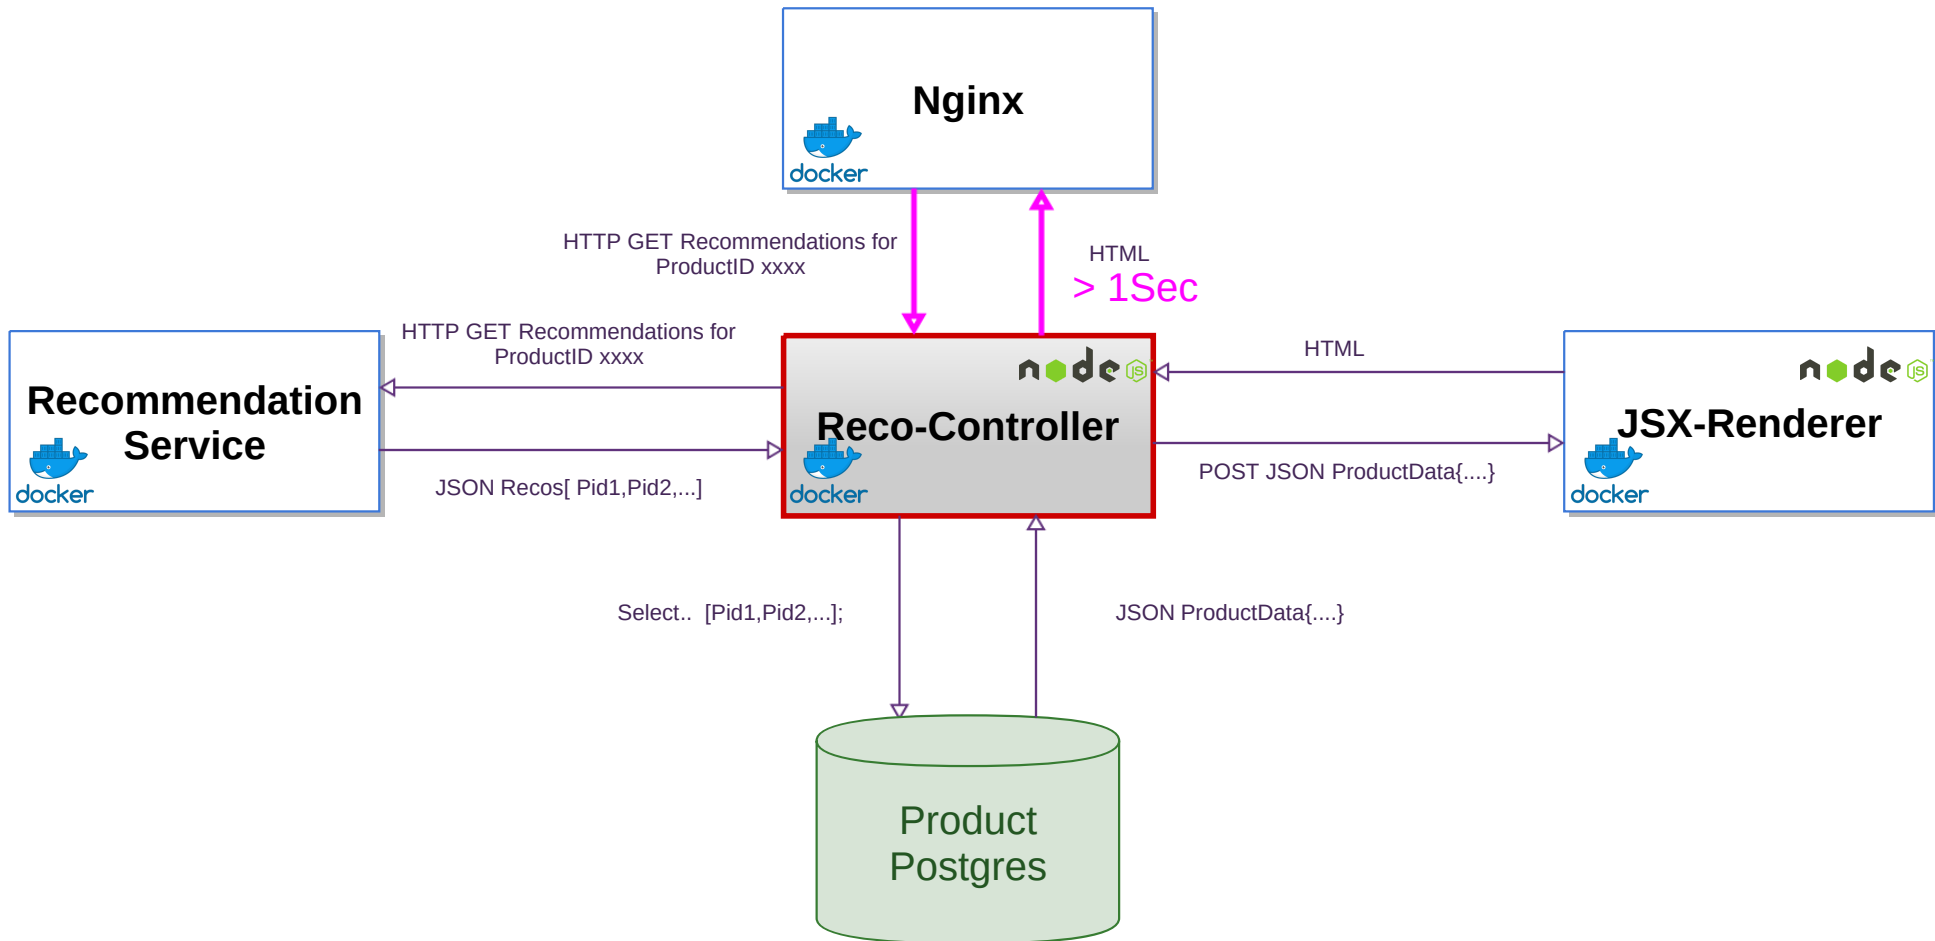












STEP 1

STEP 1

- Steady and reproducible environment

STEP 1

- Steady and reproducible environment
- Documentation of every step: idea, change, impact

STEP 1


- Steady and reproducible environment
- Documentation of every step: idea, change, impact
- Periodic and automatic tests to reproduce your issue

STEP 1

- Steady and reproducible environment
- Documentation of every step: idea, change, impact
- Periodic and automatic tests to reproduce your issue
- I use Gatling

STEP 1

- Steady and reproducible environment
- Documentation of every step: idea, change, impact
- Periodic and automatic tests to reproduce your issue
- I use Gatling



STEP 2

STEP 2

- Add monitoring to your NodeJS application

STEP 2

- Add monitoring to your NodeJS application
- System which saves and displays your monitoring data

STEP 2

- Add monitoring to your NodeJS application
- System which saves and displays your monitoring data
- Prometheus + Grafana

STEP 2

- Add monitoring to your NodeJS application
- System which saves and displays your monitoring data
- Prometheus + Grafana
- Don't use average! Use percentiles like 95p to see spikes!



STEP 2

- Add monitoring to your NodeJS application
- System which saves and displays your monitoring data
- Prometheus + Grafana
- Don't use average! Use percentiles like 95p to see spikes!
- `npm i --save express-prom-bundle 5.1.0`

STEP 2

- Add monitoring to your NodeJS application
- System which saves and displays your monitoring data
- Prometheus + Grafana
- Don't use average! Use percentiles like 95p to see spikes!
- `npm i --save express-prom-bundle 5.1.0`

```
const metricsMiddleware = require('express-prom-bundle')({
  percentiles: [ 0.5, 0.75, 0.9, 0.95, 0.99],
  metricType: 'summary',
  maxAgeSeconds: 300,
  ageBuckets: 5
})
const app = express()
app.use(metricsMiddleware)
```



STEP 3

STEP 3

- Add more monitoring

STEP 3

- Add more monitoring
- Monitor critical code paths

STEP 3

- Add more monitoring
- Monitor critical code paths
- Calls to other Services



STEP 3

- Add more monitoring
- Monitor critical code paths
- Calls to other Services
- CPU intensive paths, e.g. json-parsing or crypto


```
const promClient = require('prom-client')
const externalCallMonitor = new promClient.Summary({
  name: 'external_call_duration_ms',
  percentiles: [0.5, 0.75, 0.9, 0.95, 0.99],
  maxAgeSeconds: 300,
  ageBuckets: 5
})
```

```
const promClient = require('prom-client')
const externalCallMonitor = new promClient.Summary({
  name: 'external_call_duration_ms',
  percentiles: [0.5, 0.75, 0.9, 0.95, 0.99],
  maxAgeSeconds: 300,
  ageBuckets: 5
})
```

```
const requeststart = process.hrtime()
request({uri: "http://reco-server/reco?pid=1"}).
then((body) => {
  const requestduration = process.hrtime(requeststart)
  externalCallMonitor.observe(requestduration[1] / 1000000)
})
```



JSXrenderCall in ms



STEP 4

STEP 4

- Add more monitoring

STEP 4

- Add more monitoring
- Monitor the whole business code path


STEP 4

- Add more monitoring
- Monitor the whole business code path
- From request handling to writing the response

```
const promClient = require('prom-client')
const businessLogicMonitor = new promClient.Summary({
  name: 'business_logic_duration_ms',
  percentiles: [0.5, 0.75, 0.9, 0.95, 0.99],
  maxAgeSeconds: 300,
  ageBuckets: 5
})
```

```
const promClient = require('prom-client')
const businessLogicMonitor = new promClient.Summary({
  name: 'business_logic_duration_ms',
  percentiles: [0.5, 0.75, 0.9, 0.95, 0.99],
  maxAgeSeconds: 300,
  ageBuckets: 5
})
```

```
router.get('/recos', (req, res) => {
  const businessLogicStart = process.hrtime()
  businessLogic(req.body) // Transform Data, Call 3 external Services
    .then((html) => {
      const businessDuration = process.hrtime(businessLogicStart)
      businessLogicMonitor.observe(businessDuration[1] / 1000000)
      res.send(html)
    })
})
```



STEP 5

STEP 5

- CPU Profiling

STEP 5

- CPU Profiling
- `node --prof app.js` | Was not an option

STEP 5

- CPU Profiling
- `node --prof app.js` | Was not an option
- `npm i --save v8-profiler-node8`

```
const express = require('express')
const router = express.Router()
const profiler = require("v8-profiler-node8")

router.get('/cpuprofile', (req, res) => {
  const id = Date.now() + ".profile"
  profiler.startProfiling(id)
  // stop profiling in n seconds and exit
  setTimeout(() => {
    res.set('Content-Type', 'application/json-home')
    res.json(profiler.stopProfiling(id))
    profiler.deleteAllProfiles()
  }, 10000)
})
module.exports = router
```

Chrome analysis

STEP 6

STEP 6

- Heap Dumps

STEP 6

- Heap Dumps
- `npm i --save heapdump`


```
const express = require('express')
const router = express.Router()
const heapdump = require('heapdump')
const fileSystem = require('fs')
router.get('/heapdump', (req, response) => {
  heapdump.writeSnapshot(function(err, filename) {
    var stat = fileSystem.statSync(filename);
    response.writeHead(200, {
      'Content-Type': 'application/octet-stream',
      'Content-Length': stat.size
    });
    var readStream = fileSystem.createReadStream(filename);
    readStream.pipe(response);
  });
})
module.exports = router
```

Chrome analysis

STEP 7

STEP 7

- Add more monitoring

STEP 7

- Add more monitoring
- Monitor GarbageCollection and Internals of nodes

STEP 7

- Add more monitoring
- Monitor GarbageCollection and Internals of nodes
- `npm i --save prometheus-gc-stats`

```
const prometheus = require('prom-client')
prometheus.collectDefaultMetrics()

const gcStats = require('prometheus-gc-stats')
const startGcStats = gcStats(prometheus.register)
startGcStats();
```

Demo Board

STEP 8

STEP 8

- Tracing with DataDog

STEP 8

- Tracing with DataDog
- `npm i --save dd-trace`

STEP 8

- Tracing with DataDog
- `npm i --save dd-trace`

```
const tracer = require('dd-trace').init()
// enable and configure postgresql integration
tracer.use('pg', {
  service: 'pg-cluster'
})
// enable and configure express integration
tracer.use('express', {
  service: 'express-cluster'
})
```

DD-DEMO

MY NEXT STEPS

MY NEXT STEPS

- Local setup to reproduce the issue: docker-compose + blockade

MY NEXT STEPS

- Local setup to reproduce the issue: docker-compose + blockade
- Take a deeper look into tcp.connect, [http|tcp]_keep_alive

MY NEXT STEPS

- Local setup to reproduce the issue: docker-compose + blockade
- Take a deeper look into tcp.connect, [http|tcp]_keep_alive
- Rewrite in scala: service issue or environment issue?

QUESTIONS ?

IDEAS ?

Slides: <https://github.com/inoio/node-perf-talk>



@caraboides